

```
%matplotlib inline

from pathlib import Path

import numpy as np
import pandas as pd
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
from sklearn.metrics import roc_curve, auc


!pip install mord
from mord import LogisticIT
import matplotlib.pyplot as plt
import seaborn as sns

!pip install dmba
from dmba import classificationSummary, gainsChart, liftChart
from dmba.metric import AIC_score
```

```
⇒ Requirement already satisfied: mord in /usr/local/lib/python3.11/dist-packages (0.7)
Requirement already satisfied: dmba in /usr/local/lib/python3.11/dist-packages (0.2.4)
Requirement already satisfied: graphviz in /usr/local/lib/python3.11/dist-packages (from dmba) (0.20.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from dmba) (3.10.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from dmba) (1.26.4)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (from dmba) (2.2.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (from dmba) (1.6.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from dmba) (1.13.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (1.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (4)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (24.2)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba) (3.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib->dmba)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas->dmba) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas->dmba) (2025.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn->dmba) (1.4)
```


Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn->dmdb)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib)

```
mowers = pd.read_csv("RidingMowers.csv")  
mowers.head()
```



	Income	Lot_Size	Ownership
0	60.0	18.4	Owner
1	85.5	16.8	Owner
2	64.8	21.6	Owner
3	61.5	20.8	Owner
4	87.0	23.6	Owner

1. What percentage of households in the study were owners of a riding mower?




```
num_owners = mowers[mowers["Ownership"] == "Owner"].shape[0]  
num_owners
```

 12

```
total_households = mowers.shape[0]  
total_households
```

 24

```
percent_owners=(num_owners/total_households) * 100  
percent_owners
```

 50.0

The percentage of households in the study that were owners is 50%

2. Use all the data to fit a logistic regression of ownership on the two predictors. Remember to create dummy variables, if appropriate.

```
mowers.info()
```

```
>>> <class 'pandas.core.frame.DataFrame'>
RangeIndex: 24 entries, 0 to 23
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Income      24 non-null    float64
 1   Lot_Size    24 non-null    float64
 2   Ownership   24 non-null    object
dtypes: float64(2), object(1)
memory usage: 708.0+ bytes
```

```
mowers.Ownership.unique()
```

```
>>> array(['Owner', 'Nonowner'], dtype=object)
```

Changed ownership into category

```
mowers['Ownership'] = mowers['Ownership'].astype('category')
```

Created Dummies

```
mowers = pd.get_dummies(mowers, columns=['Ownership'], prefix='Ownership', drop_first=True)
```

```
print(mowers.columns)
```

```
Index(['Income', 'Lot_Size', 'Ownership_Owner'], dtype='object')
```

```
X = mowers[['Income', 'Lot_Size']]
y = mowers['Ownership_Owner']
train_X, valid_X, train_y, valid_y = train_test_split(X, y, test_size=0.4, random_state=1)
```

```
logit_reg = LogisticRegression(random_state=1, C=1, solver='liblinear')
logit_reg.fit(train_X, train_y)
```

```
LogisticRegression
LogisticRegression(C=1, random_state=1, solver='liblinear')
```

```
logit_reg.intercept_
```

```
array([-0.40424044])
```

```
logit_reg.coef_
```

```
array([[ 0.05457264, -0.16890291]])
```

```
pd.DataFrame({'coef': logit_reg.coef_[0], 'variable': X.columns})
```

```
coef variable
0  0.054573  Income
1 -0.168903  Lot_Size
```

```
odds_ratios_mowers = pd.DataFrame({'coef': logit_reg.coef_[0], 'odds': np.e**logit_reg.coef_[0], 'variable': X.columns})
print(odds_ratios_mowers)
```

```
coef odds variable
0  0.054573  1.056089  Income
```

1 -0.168903 0.844591 Lot\_Size

### 3. Generate the odds ratios and interpret insights.

Income (Odds Ratio = 1.0561) A 1-unit increase in income increases the odds of owning a riding mower by 5.61%.

$((1.056089 - 1) \times 100 = 5.61\%)$  An odds ratio  $> 1$ , means higher income increases ownership probability.

Lot Size (Odds Ratio = 0.8449) A 1-unit increase in lot size decreases the odds of owning a riding mower by 15.51%.

$((0.8446 - 1) \times 100 = -15.54\%)$  An odds ratio  $< 1$ , means larger lot sizes decrease ownership probability.

Start coding or [generate](#) with AI.

```
logit_reg_prob = logit_reg.predict_proba(valid_X)
logit_reg_pred = logit_reg.predict(valid_X)
```

```
logit_result = pd.DataFrame({'actual': valid_y,
                             'p_0': [p[0] for p in logit_reg_prob],
                             'p_1': [p[1] for p in logit_reg_prob],
                             'predicted': logit_reg_pred})
logit_result
```



	actual	p_0	p_1	predicted
13	False	0.738075	0.261925	False
18	False	0.466319	0.533681	True
3	True	0.636731	0.363269	False
14	False	0.443508	0.556492	True
20	False	0.642792	0.357208	False
17	False	0.666402	0.333598	False
10	True	0.791980	0.208020	False
4	True	0.411572	0.588428	True
2	True	0.626267	0.373733	False
19	False	0.477580	0.522420	True

```
classes = ['Owner', 'Nonowner']  
print(classificationSummary(valid_y,logit_reg.predict(valid_X)))  
print(classificationSummary(train_y,logit_reg.predict(train_X)))
```



Confusion Matrix (Accuracy 0.4000)

```
      Prediction  
Actual 0 1  
      0 3 3  
      1 3 1
```

None

Confusion Matrix (Accuracy 0.5714)

```
      Prediction  
Actual 0 1  
      0 2 4  
      1 2 6
```

None

Double-click (or enter) to edit

4. Among nonowners, what is the percentage of households classified correctly?

From the confusion matrix(valid),

Correctly Classified Nonowners

True Negatives (TN): 3 (Nonowners correctly classified as Nonowners).

False Positives (FP): 3 (Nonowners incorrectly classified as Owners).

◆ Percentage of Nonowners Correctly Classified:

True negative/total Nonowners

true negative = 3

Total non owners= 6

% of non owners=  $3/6 \times 100 = 50\%$

50.00% of household was classified correctly among non owners.

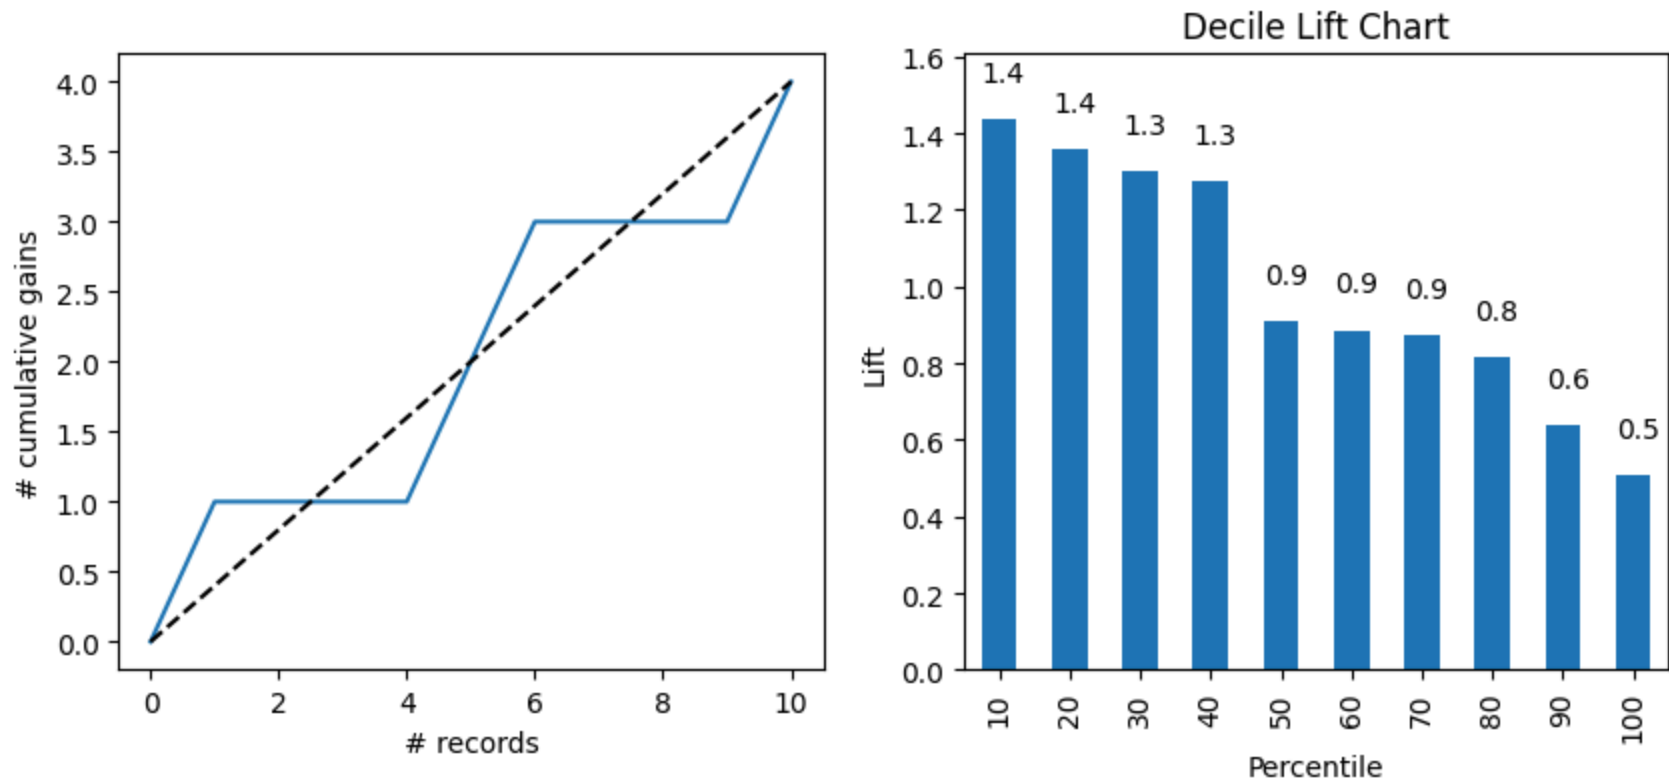
```
logit_result = logit_result.sort_values(by=['p_1'], ascending=False)
```

```
logit_result.info()
```

```
➦ <class 'pandas.core.frame.DataFrame'>
Index: 10 entries, 4 to 10
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   actual      10 non-null    bool
1   p_0         10 non-null    float64
2   p_1         10 non-null    float64
3   predicted   10 non-null    bool
dtypes: bool(2), float64(2)
memory usage: 260.0 bytes
```

```
fig, axes = plt.subplots(1, 2, figsize=(10,4))
gainsChart(logit_result.actual, ax=axes[0])
liftChart(logit_result['p_1'], ax=axes[1])
```


↔ <Axes: title={'center': 'Decile Lift Chart'}, xlabel='Percentile', ylabel='Lift'>



```
logit_result = logit_result.sort_values(by='p_1', ascending=False)
gainsChart(logit_result.actual)
```




↩➤ <Axes: xlabel='# records', ylabel='# cumulative gains'>



```
fpr, tpr, _ = roc_curve(logit_result['actual'], logit_result['p_1'])
roc_auc = auc(fpr, tpr)
plt.figure(figsize=[5, 5])
plt.plot(fpr, tpr, color='purple',
         lw=2, label='ROC curve (area = %0.4f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='blue', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.legend(loc="lower right")
```

 <matplotlib.legend.Legend at 0x7c009e7ca090>



Double-click (or enter) to edit

5. What is the probability that a household with a \$60K income and a lot size of 20,000 ft<sup>2</sup> is an owner?

```
Ownership_Nonowner_Riding_Mower= pd.DataFrame({'Income':[60.0], 'Lot_Size':[20.0]})
Ownership_Nonowner_Riding_Mower
```



```
logit_reg.predict_proba(Ownership_Nonowner_Riding_Mower)
```



```
array([[0.6243235, 0.3756765]])
```

Double-click (or enter) to edit

```
#From the array displayed,  
#The probability of being a non-owner is 0.6243 (62.43%).  
#The probability OF a household being an owner is 0.3757 (37.57%).
```

6. What is the classification of a household with a \$60K income and a lot size of 20,000 ft<sup>2</sup>? Use cutoff = 0.5

```
logit_reg.predict(Ownership_Nonowner_Riding_Mower)
```



```
array([False])
```